

AUTHBC: Feasibility Boundaries for Authenticated UAV Telemetry — An Exclusion Bound and a Capacity Envelope

Mohamed A. Farouk

[AFFILIATION – TO BE COMPLETED]

[Department, Institution, City, Country]

mohamed.besco1998@gmail.com

Abstract—A UAV that broadcasts signed, hash-chained telemetry spends more bytes on authentication than on data: a 64 B signature and a 32 B chain link accompany a 24.0 B record. AUTHBC makes the frame the unit of design: a format, a sender and a receiver in which one signature and one chain link cover b records and every frame decodes and verifies alone. Batch size, encoding and signature scheme are chosen together under freshness, verifiability, channel capacity and receiver CPU. Two boundaries follow. Where the frame does not fit: when every frame must verify alone, eight of the twelve EU863-870 LoRaWAN data rates cannot carry one. Where it does: at 50 records/s and a 100 ms deadline the design sends 43.25 B per record against 145.77 B for per-record signing, and in ns-3 one 802.11a collision domain delivers 95% of frames among 124 nodes instead of 35. Batching and header design decide what can run; among 64 B elliptic-curve signatures the scheme does not. Sender energy is measured on a Raspberry Pi 4; contention is simulated; record sizes are synthetic, checked against twelve public flight logs and an autopilot in simulation.

Index Terms—UAV telemetry, broadcast authentication, digital signatures, hash chain, batching, IEEE 802.11 DCF, LoRaWAN, ns-3.

I. INTRODUCTION

Unmanned aircraft in a swarm tell each other where they are many times a second. If those reports are to be evidence afterwards — of who was where, and of what each vehicle claimed — each one must be attributable to its sender and the sequence must be tamper-evident. The textbook construction signs every record and links it to its predecessor by a hash. The textbook cost is that the protection is larger than the thing protected: an Ed25519 or ECDSA signature is 64 B and a SHA-256 link is 32 B, while a compactly encoded position, velocity and status record is about 24.0 B. On a shared radio channel those bytes are airtime, and airtime decides how many neighbours can be heard.

Practice today signs per message. MAVLink 2 appends a 13 B symmetric tag to each message [1]; vehicular standards attach an ECDSA signature, and periodically a certificate, to each one [2]. The obvious alternative is to let several records share one frame, one signature and one chain link. Batching is not new, and its price is known in kind: records wait for the frame to fill, and a lost frame loses all of them. What a designer lacks is the accounting. What must a frame carry so that a receiver can verify it without any other frame? On

which links does such a frame not fit, whatever the batch? Where it fits, which constraint caps the batch, and what do the saved bytes buy on a contended channel, in nodes and not in percent?

We answer these for one concrete system. AUTHBC — authenticated telemetry for a blockchain-style, hash-chained ledger — is a frame format, a sender and a receiver. We built it, measured it, and searched its parameters. Our contributions are:

- **A frame as one object.** The design and its baseline are frames that were emitted, decoded and verified, and their sizes are measured lengths (Sec. IV); a size that is a sum of measured parts is marked as one. We give the field-level byte budget of two formats: a first, text-keyed one and a lean one that is 42% smaller.
- **An exclusion bound with its scope.** On LoRaWAN, eight of twelve EU863-870 data rates cannot carry one signed, chained frame; we state the four conditions under which that holds and what each relaxation buys (Sec. VII).
- **A ladder that separates causes.** Sharing a frame, sharing a signature and delta coding are separate steps, and we report bytes, capacity, receiver CPU and energy for each (Sec. VIII). Capacity is found by simulating each configuration at its own frame size and rate, with 30 seeds and a bootstrap interval, not by applying one load ceiling to all.
- **Constraints that bind in an unexpected order.** Freshness, not the MTU, caps the batch at four records. A delta-coded frame that depends on its predecessor misses the verifiability target at its own loss budget, so every frame carries one self-contained record. One CPU core keeps up with a neighbourhood of 296 nodes with Ed25519 and of 10 with BLS.

The message is short: batching and header design decide what can run, and among standard 64 B elliptic-curve signatures the scheme does not change it. AUTHBC introduces no new cryptographic primitive. What is measured, what is simulated and what is only modelled is stated with each result and collected in Sec. X.

II. BACKGROUND AND RELATED WORK

Per-message authentication. MAVLink 2 signing appends a link identifier, a 48-bit timestamp and a 48-bit truncated SHA-256 tag computed with a shared secret [1]. It authenticates a *link*: any holder of the key can produce any message, so a record cannot be attributed to one vehicle afterwards. Rotta and Mykytyn likewise attach an HMAC and a timestamp to each broadcast telemetry message of a swarm [3]. Vehicular broadcast uses public-key signatures: ECDSA over each message, with the signer’s certificate in some messages and an eight-octet digest of it in the others [2], [4]. That is the trust model we adopt, and the per-message cost we start from.

Signing streams. To amortise one signature over many packets, Gennaro and Rohatgi sign the first packet and chain the rest by hashes [5]; one lost packet breaks the chain. Wong and Lam sign a hash tree over a block and send in every packet the hashes needed to reach the root [6]: each packet verifies alone, and the signing cost is amortised but the bytes are not. EMSS tolerates loss by carrying several hashes per packet and a periodic signature packet, and TESLA replaces the signature by a MAC whose key is disclosed later [7], [8]; TBRD applies TESLA to UAS Remote ID and reports about half the overhead of signatures [9]. These schemes spread *one authenticator across packets*, so packets depend on each other or on a later one. We put *several records in one packet*, which a small telemetry record allows: each frame stays self-verifying and the cost moves to latency, at the sender. TESLA delays verification at the receiver instead, and a disclosed MAC key gives no non-repudiation, which a ledger needs. Section IX places these schemes at our operating point.

Aggregate and batch verification. BLS signatures aggregate [10], [11], and certificateless aggregate schemes (CLAS) for vehicular networks build on such constructions [12]. Zhang *et al.* verify many signatures at once at a roadside unit [13], and Rajasekaran *et al.* batch-authenticate a group of drones in a pairing-based scheme, reporting a communication cost that grows linearly with the group [14]. All of these reduce the verifier’s *work*, and an aggregator may forward one short aggregate in place of many signatures; but each message, as its signer sends it, still carries its own signature material (Sec. IX).

Hash-linked records. Linking each record to the hash of its predecessor to make a sequence tamper-evident is due to Haber and Stornetta [15]; blockchain uses on UAVs are surveyed in [16]. We take the chain as given and ask what it costs on air.

Small payloads. That a signature leaves little room in a small frame is known: Gündoğan *et al.* find 9 B of application data left beside an Ed25519 signature in an 802.15.4 frame [17], and post-quantum signatures do not fit one 5G broadcast block [18]. For Remote ID, a 64-octet signature is pagged over 25-octet Bluetooth frames with a parity page [19]: the spreading and recovery that our scope excludes. Section VII applies the same arithmetic to LoRaWAN with the chain link included. Compact encodings [20] and delta coding of time series [21] are established; a lossy link limits how far

Per-record signing: one frame per record, 145.8 B each

H_f	link	key	signature	$\times 4$
-------	------	-----	-----------	------------

AUTHBC: one frame per four records, 173.0 B

H_f	link	key	δ	δ	δ	signature
23	35	24.0	3×9.0			64

Fig. 1. The lean frame, to scale (bytes below). Per-record signing repeats header, chain link and signature with every record; the design carries them once per four records, and the three records after the first are coded as differences (δ).

a delta chain may run (Sec. VIII-B).

Channel models. Broadcast frames are not acknowledged and their contention window never doubles, so Bianchi’s saturation model of the 802.11 DCF [22] does not apply as it stands; we use the broadcast model of Ma and Chen [23], [24] and validate it against ns-3 [25]. For LoRa, Haxhibeqiri *et al.* model single-gateway scalability [26] and Zirak *et al.* report air-to-air delivery against range on hardware [27].

III. SYSTEM AND THREAT MODEL

Network and traffic. N UAVs share one channel and hear each other: one collision domain, which we call a neighbourhood of N nodes; each node hears the other $N - 1$. Each produces Λ telemetry records per second and broadcasts them; there are no acknowledgements and no retransmissions. A record holds a millisecond timestamp, latitude, longitude and altitude, three velocity components, battery level and flight mode, all integers. Our primary operating point is $\Lambda = 50$ records/s with a freshness bound $D_{\max} = 100$ ms: the rate is PX4’s companion-computer stream rate [28] and the pair meets the direct UAV-to-UAV broadcast requirements of 3GPP [29]. A relaxed point, 20 records/s and 250 ms, is reported for comparison.

Ledger. Each UAV keeps an append-only chain of its own records, $prev_hash_i = \text{SHA-256}(rec_{i-1})$, and stores the verified records it hears from others. There is no consensus layer; the subject of this paper is authenticated dissemination.

Adversary. The adversary controls the radio: it reads, injects, replays, modifies, delays and drops frames, and may operate many radios. It holds no honest UAV’s private key.

What is protected. (i) *Origin and integrity*: a stored record was produced by the UAV it names and was not altered, by a public-key signature checked before anything is stored. (ii) *Replay*: sequence numbers are strictly increasing per sender; a repeated frame is refused. (iii) *Equivocation*: two validly signed records for one sequence number are kept as a pair, which is transferable evidence because both carry the sender’s signature. (iv) *Continuity*: omission or reordering within what a receiver holds is detected through the chain.

What is not. Telemetry is sent in clear. Jamming is not resisted: loss appears as the robustness measure V below, not as a security property. Replay is refused by sequence number, not by time, so a frame that an attacker withholds and delivers later is accepted, authentic and in order, however

TABLE I
BYTE BUDGET OF A FRAME BY FIELD: FIVE-DIGIT SENDER IDENTIFIER,
ONE HOUR INTO A 50 HZ FLIGHT (SENDER 0, FIRST FRAME OF A FLIGHT,
IN PARENTHESES).

field	content	first	lean
(map)	CBOR map header	1	1
v, t	version; frame type	6	4
src	sender identifier, 16 bit	7 (5)	4 (2)
base_seq	sequence no. of first record	14 (10)	6 (2)
n	records in the frame	3	2
recs	key, length of records	6	3
auth	key, length of signature	7	3
header H_f		44 (38)	23 (17)
chain link		32 per record	35 per frame
signature g_a (Ed25519, ECDSA)		64	64
keyframe record s_k		59.9*	24.0
delta record s_δ , 50 ms apart		44.0*	9.0

* includes that record's own 32 B chain link.

stale; refusing it needs a clock shared with the sender, which we do not assume. A UAV with a valid key can sign false data; only inconsistent falsehood is detectable. Key distribution, revocation and admission are assumed: receivers hold senders' public keys, and we charge the bytes of delivering certificates separately (Sec. VIII-E). Senders are not anonymous.

Verifiability. V is the probability that a record a UAV generates is received, decoded and verified by a given neighbour. The requirement is $V \geq 1 - \epsilon$ with $\epsilon = 0.05$.

IV. THE AUTHBC FRAME

A frame is one canonical CBOR map [20] with integer keys (Table I). `src` and `base_seq` name the sender and the sequence number of the first record; `link` is the hash of the record before it; `recs` is the record stream and `auth` one signature. A record in the stream is nine variable-length integers. The first record of a frame is a *keyframe*, carrying absolute values (24.0 B on average); each later one is a *delta* from the record before it (9.0 B when records are 50 ms apart). Sender and sequence number are not repeated in the records.

What is signed is not the stream but the canonical ledger records, each with its own `prev_hash`. The receiver rebuilds them — the first `prev_hash` is `link`, each later one is the hash of the record just rebuilt — and only then checks the signature. A frame that cannot be decoded therefore cannot be verified, and nothing unverified is stored. Each verified record is offered to the ledger store, which refuses a sequence number that is not newer, keeps an equivocating pair as evidence, and flags a link that contradicts a record it holds. The state against which a later frame would be decoded advances only after a frame has been verified and stored, so a replayed frame cannot desynchronise a stream. Flipping any single bit of a frame makes verification fail (all 1392 positions of a four-record frame were tried).

When a frame is lost, the next one still decodes and verifies, because it starts with a keyframe. Its `link` then names a record this receiver does not hold: the gap shows in the sequence numbers, and the link is stored unchecked,

to be checked by any party that holds the missing record. A late copy of the lost frame is refused, since sequence numbers must increase; the gap is not repaired.

One link per frame. Inside a frame the order and content of records rest on the frame signature alone; in the first format below each record also carries an independent hash. A frame is signed and accepted as a unit, so no strength is lost, but one mechanism remains where there were two.

The first format. For comparison we keep the format we built first: text keys, and in every record its own chain link, sender identifier and sequence number. Its header is 44 B, of which 29 B are key names. That format is implemented and frozen. Its sizes in this paper are sums of separately measured parts, and its delta-coded frames were never emitted; we mark such sizes wherever they appear.

Design parameters. A configuration is the batch size b ; whether a signature covers each record (*inline*; `auth` is then an array of b signatures) or the frame (*batch*); the signature scheme (Ed25519 [30] and ECDSA P-256 [31] at 64 B, BLS12-381 [10] at 96 B in the library used); and the reference interval r : one frame in r starts with a keyframe, the others with a delta against the previous frame.

V. CONSTRAINTS

Bytes. With header H_f , link field L , signature g_a , keyframe s_k and delta s_δ , a self-contained batch frame costs, per record,

$$B(b) = \frac{H_f + L + g_a + s_k + (b-1)s_\delta}{b}. \quad (1)$$

$B(1)$ is our baseline: one record, one signature, one frame. Signing each record of a shared frame replaces g_a by $b g_a$. Equation (1) is exact to within a byte of CBOR length prefix; for the lean design and its baseline we report the lengths of emitted frames.

Freshness. The oldest record in a frame waits for the frame to fill. Adding the channel-access and airtime delay T_{ch} and the verification time t_v ,

$$D(b) = \frac{b}{\Lambda} + T_{ch} + t_v \leq D_{\max}. \quad (2)$$

Fill time dominates. At both operating points $\Lambda D_{\max} = 5$: five records fill in exactly $D_{\max}$ and leave nothing for the channel, so $b \leq 4$ whatever the encoding or the scheme. The 1500 B MTU would admit 31 records in the first format and 151 in the lean one.

Verifiability. If one frame in r is self-contained, frame j of a group verifies only if it and the $j-1$ frames before it arrive. With independent frame loss p ,

$$V(r, p) = \frac{1}{r} \sum_{j=1}^r (1-p)^j, \quad (3)$$

which equals $1-p$ at $r=1$ and is smaller for every $r>1$. A target $V \geq 1-\epsilon$ with $\epsilon \leq p$ therefore admits $r=1$ only. The same argument forbids splitting a signed unit over n frames, which verifies with probability $(1-p)^n$.

Channel. For a configuration — a frame size and a frame rate per node — $N_{\max}$ is the largest neighbourhood in which

the mean delivered fraction of frames is at least $1 - \epsilon$. We find it by simulation (Sec. VI). As a conservative companion we give N_{sat} , the largest N whose offered frame rate does not exceed the saturation throughput of the broadcast DCF model [24].

Receiver CPU. A receiver verifies every neighbour's frames. With k signatures per frame, verification time t_v and hash time t_h , one neighbour costs $(\Lambda/b) k t_v + \Lambda t_h$ seconds of CPU per second, and one core serves

$$N_{\text{cpu}} = 1 + \left\lfloor \frac{1}{(\Lambda/b) k t_v + \Lambda t_h} \right\rfloor \quad (4)$$

nodes.

Payload. A frame that verifies alone must fit the link's maximum payload M :

$$M \geq H_f + L + g_a + s_k. \quad (5)$$

The residual room beside a signature has been computed before for 802.15.4 [17]; here the chain link and a record that decodes alone are charged as well.

Relayed records. A node that forwards records of b different originators cannot re-sign them. It carries their b signatures, $64 b$ B, or one BLS aggregate of 96 B [11], saving $64 - 96/b$ B per record at the price of a pairing-based verification. Both placements are implemented. We evaluate only a sender's own records; multi-hop forwarding is outside this paper.

VI. METHOD

Implementation. Formats, sender, receiver and store are one Python code base with four record encodings (JSON, CBOR, MessagePack, delta) and three signature schemes behind common interfaces, each scheme checked against published test vectors (for Ed25519, those of [32]). Every table in this paper is generated from result files that the code writes; a gate re-derives the deterministic ones on each commit.

Telemetry. Records come from a seeded random-walk generator: 30 seeds of 1000 records at 50 ms spacing, that of the relaxed point. At the 20 ms of the primary point its delta record is the same size, the floor of the format. Header fields are sized for a five-digit sender identifier one hour into a 50 Hz flight. Section VIII-G checks the generator against flight logs.

Timing and energy. Signing, verification, hashing and encoding are timed on a Raspberry Pi 4 (medians; Ed25519 sign/verify 88/259 μ s, ECDSA 136/327 μ s, BLS 2.73/8.61 ms, SHA-256 2.75 μ s). Sender energy is metered on the same board with an INA219 on the supply rail, five repetitions per configuration, idle power subtracted.

Channel simulation. ns-3.48 [25]: IEEE 802.11a at 6 Mb/s, ad hoc mode, broadcast frames, all nodes in range at equal received power, so there is no capture and no hidden terminal. Each node sends one frame every b/Λ seconds. Its send instant is drawn uniformly within each period, so the relative phases of the nodes change from frame to frame. A strictly periodic source freezes each node's phase for a whole run. A run then samples a single phase configuration, in which a few nodes may collide in every period while the others lose nothing. The




Fig. 2. Payload limits of the twelve EU863-870 data rates against the smallest lean frame and the frames of this telemetry (64 B signature).

TABLE II
DOES ONE SELF-CONTAINED LEAN FRAME FIT? FLOOR: EVERY FIELD ZERO. RANGE: FRAMES OF THIS TELEMETRY.

	floor (B)	range (B)	50/51 B 5 rates	115 B 3 rates	222/242 B 4 rates
64 B signature, chain link	124	134–147	no	no	yes
... link not sent on air	89	99–112	no	yes	yes
48 B BLS signature	108	118–131	no	no*	yes
13 B symmetric tag	72	82–95	no	yes	yes

* a frame of zeros would fit; no frame of this telemetry does.

two sources deliver the same on average — over 14 points and 1 260 runs the means differ by 0.0004 ± 0.0009 — but the strictly periodic one varies 6–23 times more from run to run. A run lasts 20 s; each point is 30 runs. N_{max} is the largest simulated N whose mean delivered fraction is at least 0.95 with every smaller simulated N also passing; its interval is the 2.5–97.5 percentile of that quantity over 10 000 bootstrap resamples of the seeds. The saturation models were validated separately against saturated ns-3 runs (Sec. VIII-H).

VII. LOW-RATE LINKS: WHERE NO BATCH HELPS

The first question is not what a design saves but whether its frame can be sent at all.

LoRaWAN limits the payload of a frame by data rate, and its EU863-870 plan defines twelve data rates [33]: seven LoRa (DR0–DR6), one FSK (DR7) and four LR-FHSS (DR8–DR11). By Eq. (5) a frame that verifies alone needs a header, a chain link, a signature and one record that decodes alone. The smallest such frame the lean format can emit at all — sender 0, sequence 0, every field zero — is $16+35+64+9 = 124$ B, and we emitted, signed and verified it. Header, link and signature alone are 115 B.

Figure 2 sets that against the payload limits. Five data rates (DR0–DR2, DR8, DR10) allow 50 or 51 B: the signature alone does not fit. Three (DR3, DR9, DR11) allow 115 B: header, link and signature fill the payload and no record fits. Four (DR4–DR7) allow 222 or 242 B and are feasible. **Eight of twelve data rates cannot carry one signed, chained frame,**

in either format and under either of the standard’s payload tables. None of the eight depends on the telemetry, since each is excluded against the format’s floor.

Scope. The result holds under four conditions. (i) The limits are LoRaWAN’s regional parameters; the LoRa PHY itself carries 255 B [34] and is not excluded. (ii) Every frame carries a public-key signature. (iii) There is no loss recovery. (iv) $\epsilon \leq p$, which by Eq. (3) forbids spreading a signature over several frames; at $p = 0.01$ and $\epsilon = 0.05$ a signed unit may span five frames and the exclusion no longer holds as stated.

What each relaxation buys (Table II). Taking the chain link off the air frees the three 115 B rates: they are excluded by the link, not by the signature. The smallest standardised signature, 48 B BLS12-381 [35], fits a 51 B payload and leaves 3 B; at 115 B a frame of zeros would fit (108 B) but the smallest frame of our telemetry is 118 B, a three-byte miss that we do not count as an exclusion. A 13 B symmetric tag, as in MAVLink 2, fits the 115 B rates, but at 50 or 51 B header and link already take 51–58 B: a hash-chained frame does not fit there under any authenticator.

What the feasible rates carry. At DR5 under the 1% duty cycle a 242 B frame holds 8 records and sustains 0.20 records/s per node (5 records and 0.13 in the first format). The batch must be sized at the record spacing it implies: a delta that spans 5.5 s costs 13.6 B, not the 9.0 B it costs at 50 ms. The ns-3 LoRaWAN module carries at most 222 B, where a frame holds six records. Simulated with a 218 B payload every 36.4 s (0.165 records/s per node), on one channel with one demodulator and no capture, the largest neighbourhood with mean $V \geq 0.95$ is three nodes (30 seeds, 95% interval [2, 3]; one if 95% of runs must each reach it); the closed-form model of Haxhibeqiri *et al.* [26], fitted to hardware interference measurements and evaluated at the same operating point, gives four. Such a neighbourhood carries about 0.5 records/s in all, where the 802.11 design of Section VIII carries some 6 200. On these links a ledger is a slow record of provenance, not a smaller version of the same design.

VIII. RESULTS ON 802.11

On a link that carries the frame, the questions are what each design choice buys and how many nodes one channel then serves.

A. The design ladder

Table III takes the design apart. A baseline that signs every record and sends it alone confounds two things, so rows 2 and 6 separate them: four records share a frame but each keeps its signature. In the lean format that step takes 145.77 B per record to 93.00: the header and the chain link are paid once per four records, and the three records that follow the first can be coded as deltas. Sharing the signature as well gives **43.25 B, 70.3% below the baseline**; row 7 is the same frame without delta coding, 54.50 B. In the first format the steps are sharing the frame (174.25 to 141.25 B), sharing the signature (93.25) and delta coding (74.96, 57.0% below its baseline). It




Fig. 3. Bytes per record against batch size, by where the chain link sits (delta records throughout). First-format curves are sums of measured parts; lean frames are emitted.

ends higher because each of its records carries a 32 B link of its own, which no batch size amortises.

Figure 3 shows where the link belongs. Moving the first format’s link from the record to the frame saves $32(1 - 1/b)$ B per record: 24 B at $b = 4$, which gives 51.0 B. The lean header and record coding account for the remaining 7.7 B. With a link in every record, 12 records per frame still cost 54.3 B each, more than one link per frame costs at four.

An exhaustive search over placement, batch size 1–12 and reference interval returns row 8 as the byte-minimal configuration that meets the freshness, verifiability and MTU constraints, and row 4 for the first format. The batch is set by freshness (Eq. 2), not by the MTU.

The choices are made together because they are not independent. One signature per frame saves $g_a(1 - 1/b)$ per record, and delta coding saves $(1 - 1/b)$ of the difference between a self-contained record and a delta, since each frame must start with a self-contained one. Both vanish at $b = 1$, where rows 1 and 5 already are the design. Batching is the choice the other two gains depend on, and the deadline sets it.

Signature scheme does not move these rows among the 64 B schemes: Ed25519 and ECDSA produce identical frames. BLS adds 32 B per frame and, as shown below, fails on CPU.

B. Loss, and why every frame carries a keyframe

Delta coding is usually run with a keyframe every K records. With $K = 16$ and $b = 4$ only one frame in four decodes without its predecessor, and a lost frame takes the following ones with it. Table IV measures this with the codec in the loop — 7 200 frames per cell through the sender, an independent loss process and the receiver: at the loss rate the system is specified for, $r = 4$ verifies 0.876 of records, not 0.95. Equation (3) predicts the measurement at every r . Bursty loss is kinder to a dependent design — a Gilbert process with the same mean and bursts of four frames gives 0.929 at $r = 4$ — but never reaches $1 - p$. A cleaner link would admit a longer interval ($r = 4$ at $p = 0.02$, $r = 8$ at $p = 0.01$), and the whole

TABLE III

THE DESIGN LADDER AT 50 RECORDS/S AND $D_{\max} = 100$ MS. N_{sat} : SATURATION BOUND. $N_{\max}$: LARGEST NEIGHBOURHOOD WITH MEAN DELIVERY ≥ 0.95 IN NS-3 [95% INTERVAL]. CPU: SHARE OF ONE RECEIVER CORE SPENT VERIFYING AND HASHING AT $N_{\max}$.

configuration	b	frame (B)	B/record	with cert.	N_{sat}	$N_{\max}$	CPU
<i>First format: text keys, a chain link in every record[†]</i>							
1 sign every record, one record per frame (baseline)	1	174.3	174.25	213.05	18	32 [32, 32]	41%
2 sign every record, share the frame	4	565.0	141.25	150.95	24	56 [56, 56]	72%
3 one signature per frame	4	373.0	93.25	102.95	30	76 [75, 76]	25%
4 ... and delta records	4	299.9	74.96	84.66	34	88 [88, 89]	29%
<i>Lean format: integer keys, one chain link per frame</i>							
5 sign every record, one record per frame (baseline)	1	145.8	145.77	184.57	20	35 [35, 35]	45%
6 sign every record, share the frame (delta records)	4	372.0	93.00	102.70	30	76 [76, 77]	98%
7 one signature per frame, no delta coding [†]	4	218.0	54.50	64.20	42	108 [107, 109]	36%
8 one signature per frame, delta records (the design)	4	173.0	43.25	52.95	52	124 [124, 125]	42%

[†] size is a sum of measured parts; rows 5, 6 and 8 are emitted frames.

TABLE IV

VERIFIABILITY AT FRAME LOSS $p = 0.05$ WITH ONE SELF-CONTAINED FRAME IN r ($b = 4$, LEAN FORMAT; 30 SEEDS, 95% INTERVAL).

r	B/record	V measured	Eq. (3)	$V \geq 0.95$
1	43.25	0.952 [0.947, 0.957]	0.950	yes
2	41.38	0.921 [0.916, 0.927]	0.926	no
4	40.44	0.876 [0.863, 0.889]	0.881	no
16	39.73	0.655 [0.618, 0.690]	0.665	no

TABLE V

DELAY OF THE OLDEST RECORD IN A FRAME ($b = 4$, 50 RECORDS/S); CHANNEL DELAY SIMULATED AT THE HEAVIEST LOAD THAT STILL DELIVERS 95% OF FRAMES.

	fill (ms)	channel (ms)	verify (ms)	total (ms)	margin (ms)
Ed25519, mean channel delay	80	0.76	0.26	81.0	19.0
Ed25519, 99th percentile	80	2.40	0.26	82.7	17.3
Ed25519, worst observed	80	3.49	0.26	83.8	16.2
BLS, worst observed	80	3.49	8.61	92.1	7.9

prize is 3.5 B per record. The loss rate here is a specification, not a measurement: at the 2.3×10^{-4} we measured between two radios (Sec. VIII-H) every interval in the table meets the target, $r = 16$ with $V = 0.998$. The design is sized for the budget and therefore uses $r = 1$: each frame starts with a keyframe, costs 43.25 B per record, and loses only its own records when it is lost.

These figures give every frame the same loss probability. At a fixed bit error rate a longer frame is lost more often, and a batch is a longer frame: on a channel that loses 5% of the baseline's one-record frames, the design's frame is lost with probability 0.057, so 0.943 of its records are expected to verify, below the target it meets at equal frame loss; the measurement, 0.947 [0.941, 0.952], agrees with that expectation. It reaches 0.95 only while one-record frames are lost at most 0.044 of the time. The comparisons in Table III are at equal frame loss and carry that qualification.




Fig. 4. Delivered fraction of frames against neighbourhood size in ns-3: lean baseline (row 5) and design (row 8), mean and range of 30 runs.

TABLE VI

$N_{\max}$ OF BASELINE AND DESIGN AT BOTH OPERATING POINTS (NS-3, 30 SEEDS, 95% INTERVAL).

	baseline	design	ratio
<i>50 records/s, 100 ms</i>			
first format	32 [32, 32]	88 [88, 89]	2.8×
lean format	35 [35, 35]	124 [124, 125]	3.5×
<i>20 records/s, 250 ms</i>			
first format	78 [77, 78]	219 [215, 221]	2.8×
lean format	85 [85, 86]	306 [306, 308]	3.6×

C. Freshness

Table V adds the terms of Eq. (2). The frame fills in 80 ms; channel access and airtime add 0.76 ms on average and 3.49 ms at worst; one Ed25519 verification adds 0.26 ms. The oldest record is verified within 83.8 ms, leaving 16.2 ms of the 100 ms budget. A fifth record would need a fill time of 100 ms by itself. Batching b records needs a record rate above $b/D_{\max}$, here 40 records/s: at 20 records/s with the same 100 ms bound the batch is one record and nothing is gained.

D. Capacity

Bytes matter because they are airtime. Figure 4 shows delivery falling as nodes are added, and Table III gives the resulting $N_{\max}$ for every row. The lean baseline holds 35 nodes

at 95% delivery and the design 124 (3.5 $\times$); in the first format the figures are 32 and 88 (2.8 $\times$). The larger step is sending a quarter as many frames: row 6, which shares the frame but not the signature, already reaches 76 (2.2 $\times$ the baseline), and sharing the signature as well adds 1.6 $\times$. Each frame pays a preamble, an inter-frame space, a backoff and MAC framing whatever its length — about 173 μ s, 47% of what a baseline frame costs the channel — so removing three frames in four is worth more than their bytes.

These figures hold the mean over runs to the 0.95 target. Requiring that 95% of individual runs reach it gives 34 and 120 (31 and 86 in the first format), at most four nodes fewer.

Table VI repeats the comparison at the relaxed operating point. The ratio hardly moves with the operating point — 3.5 $\times$ and 3.6 $\times$ in the lean format, 2.8 $\times$ and 2.8 $\times$ in the first — and it depends on the format, so we quote it with the format and not as one number.

The capacity envelope: a rule for the crossing. At the 5% level, loss follows a line in airtime. With $f = \Lambda/b$ the frames per second a node sends and T the time one frame holds the medium (its airtime plus DIFS),

$$N_{\max} \approx \frac{0.05}{f(aT + c)}, \quad a = 0.071, \quad c = 8 \mu\text{s}, \quad (6)$$

where c is twice the simulator’s preamble-detection time and a was fitted to six rows of Table III (1, 2, 4, 5, 6 and 8), which it reproduces within 3.8%. For seven other configurations we wrote down what the rule gives before simulating them. All seven fell within 3.4% of it, and the five one-record frames of Table IX, predicted the same way, within 4.2%, all above it. The shortcut we had used before, one utilisation ceiling measured at a single point, misses the thirteen by up to 15%, in both directions. Equation (6) is a fit at one loss level and one PHY rate, checked for frames of 94–565 B, 5–50.5 frames/s and 28–306 nodes; the tables quote simulated values, not the rule.

Why the crossing is there. The slope is not a free constant. A frame is lost in this scenario in two ways. The first is known from the analysis of periodic vehicular broadcast [36]; the second is absent from that analysis. Two stations that found the medium busy may draw the same backoff slot: with occupancy $\rho \approx NfT$ and a contention window of $W = 16$ slots, counting the stations that wait together gives a loss of about $\rho^2/(W(1 - \rho))$, hence $a = \rho/(W(1 - \rho))$, which is 0.071 at $\rho = 0.53$, the mean occupancy at which loss reaches 5%. Or a station decides to send within the 4 μ s in which it cannot yet sense another: that is c . An event simulation of the access rule alone, with nothing fitted and none of the simulator’s code, reproduces all 18 simulated crossings within 1.2%. Asked for crossings where nothing had been simulated — the contention window doubled, and loss levels of 2% and 10% — it predicted all six within 1.7%. Doubling the window moves the baseline’s crossing only from 35.3 to 39.2 nodes (predicted 39.1): stations that count down for longer wait together more often. The capacities therefore follow from the access rule and one detection time.

TABLE VII
SENDER ENERGY PER RECORD ON A RASPBERRY PI 4, FIRST FORMAT, METERED (MEDIAN AND RANGE OF FIVE REPETITIONS) AGAINST THE COMPOSED MODEL.

	metered (μ J)	model (μ J)	gap
every record signed	118.8 [118.5, 119.9]	105.8	12.3%
one signature per 4	58.4 [58.1, 59.3]	54.3	7.5%

Receiver CPU. By Eq. (4) one core of the Raspberry Pi 4 serves a neighbourhood of 296 nodes when each frame carries one Ed25519 signature and of 77 when every record is signed. At its $N_{\max}$ the design uses 42% of a core, so the channel binds first. Row 6 shows the constraint is not idle: with a signature on every record it reaches the channel’s limit at 98% of a core. With BLS a verification takes 8.61 ms and one core serves 10 nodes: on this hardware BLS is excluded by CPU long before the channel is reached, in addition to its larger signature. These figures verify one signature at a time. Ed25519 was designed for batch verification — its authors report under 134 000 cycles per signature in a batch of 64, against 273 364 for one [30] — which would roughly halve the receiver’s cost; we did not use it, and ECDSA does not admit it. Equation (4) charges cryptography only. Decoding a frame and rebuilding its records is not charged, and in our interpreted prototype it costs several times the verification: these are ceilings for a compiled receiver, not what the prototype sustains.

Other PHY rates. We simulated one rate. Evaluating the saturation bound with the same model at 12 and 24 Mb/s, and on the 10 MHz channel of 802.11p at 3–12 Mb/s, gives design-to-baseline ratios of N_{sat} from 1.7 $\times$ to 4.3 $\times$, lowest at the slowest rate, because the fixed cost per frame does not shrink with the payload.

E. Certificates

The rows so far assume that receivers already hold the senders’ public keys. Vehicular practice sends a full certificate in every fifth message and an eight-octet digest in the others [4]. Charging that to every frame — a 162 B certificate, so 38.8 B per frame on average — gives the “with cert.” column of Table III: 52.95 B per record for the design against 184.57 for the baseline. Certificates are a per-frame cost, so batching amortises them too. These are explicit certificates; implicit ones are shorter (a signed message of at most 226 B against 330 B [4]) and are not priced, so the column is an upper bound.

F. Energy

Table VII meters the sender: encoding, hashing and signing a record costs 58.4 μ J in the batched configuration against 118.8 μ J when every record is signed, 51% less. The model composed from separately measured times and powers is 7.5–12.3% low, the gap being frame assembly in an interpreted prototype. These measurements are of the first format’s encoders; the lean codec has not been timed on the board. Receiver and radio energy are modelled, not metered: adding one receiver

TABLE VIII

BYTES PER MESSAGE OF CERTIFICATELESS AGGREGATE SCHEMES AS REPORTED IN [12] (WITH A 67 B MESSAGE), AND PER RECORD OF AUTHBC WITH CERTIFICATES.

	bytes per message
Wang <i>et al.</i> *	859
Liang <i>et al.</i> *	735
Xu <i>et al.</i> *	596
Cahyadi <i>et al.</i> *	583
Li <i>et al.</i> [12]	583
AUTHBC baseline (row 5), with certificates	184.57
AUTHBC design (row 8), with certificates	52.95

* total per message as reported in [12].

and the radio, the model gives 131.7 against 376.8 μJ per record.

G. Is the synthetic telemetry representative?

Every size above comes from one generator, whose delta record is exactly 9.0 B in every sample — the floor of the format. We ran twelve public PX4 flight logs (four quadrotors, two hexa/octorotors, three fixed wings, three VTOLs; 29 092 in-flight records), chosen by a rule fixed before any was opened, through the same encoder. The logs carry position at 5 Hz, so the comparison is at 0.2 s spacing and above, not at the 50 Hz of the operating point. At 0.2 s a real delta record averages 11.1 B against the generator’s 10.8 B, and a real keyframe 22.0 B against 24.0. The spread is by speed: the six multicopters, mostly hovering, average 10.0–10.5 B and the three fixed wings up to 13.2 B. With real sizes at 0.2 s the design saves 69.1% (68.2% in the least favourable log) where the generator gives 69.4%. The generator is representative to within about one percentage point, and optimistic for fast vehicles.

For the operating rate itself we flew PX4 in software-in-the-loop — its own estimator, logger and MAVLink module on a simulated quadrotor — with position logged every 20 ms. Over 14 317 in-flight records, at up to 10 m/s, every delta record is 9.0 B, the floor of the format; the keyframe averages 20.9 B and the design costs 42.47 B per record, 1.8% less than with the generator’s records. The same autopilot’s 50 Hz position stream holds a 20.0 ms period to within ± 0.4 ms (central 68%).

H. Validation of the models

The unicast DCF model [22] matches saturated ns-3 runs within +1.29/−0.40% and the broadcast model [24] within $\pm 0.51\%$ over $N = 5$ –50 (30 seeds). On two Raspberry Pi 4 radios with a single transmitter, the airtime of a 1400 B broadcast frame measures 1.995 ms against a predicted 1.988 ms (0.36%) and 4 frames in 17 600 are lost ($p = 2.3 \times 10^{-4}$). That experiment has no contention. The capacities in this paper rest on simulation.

IX. COMPARISON WITH PUBLISHED SCHEMES

Stream signing. Table IX places the schemes of Section II at our operating point. None is implemented here: each row

TABLE IX

STREAM-SIGNING SCHEMES AT 50 RECORDS/S: THE LEAN ONE-RECORD FRAME WITH EACH SCHEME’S AUTHENTICATOR, SIZED FROM ITS DEFINITION. $N_{\max}$ SIMULATED AS IN TABLE III.

scheme	auth. B/record	total B/record	N_{sat}	$N_{\max}$
MAVLink 2 tag [†]	13	93.8	23	42 [41, 42]
TESLA ^{†‡}	24	105.8	22	40 [40, 40]
Gennaro–Rohatgi [‡]	32	113.8	21	39 [39, 39]
EMSS [‡]	65	147.3	20	34 [34, 34]
Wong–Lam tree, block of 4	129	210.8	17	28 [28, 29]
signature per record (row 5)	64	145.8	20	35 [35, 35]
AUTHBC design (row 8)	16	43.3	52	124 [124, 125]

[†] no non-repudiation.

[‡] a packet does not verify alone: $N_{\max}$ is an upper bound.

is the lean one-record frame with that scheme’s authenticator, sized from its definition, in place of the signature, and $N_{\max}$ is simulated for a frame of that size. All of them send a frame per record — 50.0 frames/s per node (50.5 for EMSS, with its signature packets) where the design sends 12.5 — so the fixed channel time of a frame is paid for every record whatever the authenticator weighs. A 13 B tag in place of a 64 B signature raises $N_{\max}$ from 35 to 42 nodes and gives up non-repudiation; four records in a frame raise it to 124. Where a packet does not verify alone the simulated figure counts frames delivered, not records verified, and is an upper bound. Tree chaining is the one scheme that keeps what a ledger needs — every packet verifies alone and its origin can be shown to a third party — and it pays 129 B per packet for it. These schemes amortise the cost of signing; on a contended channel the cost that binds is the frame.

Aggregate signatures. The nearest published comparators are certificateless aggregate signature schemes for vehicular broadcast. Table VIII lists what one signer sends per message. For n messages that is n times as much, since each carries its own identity, key and signature components. Aggregation acts afterwards, at a roadside unit: it reduces the verifier’s work and what the unit forwards — in the one scheme read at source, 388 B for a single signature and 784 B for the aggregate of a hundred [37]. Batching acts before, on a sender’s own records. The two act on different links, and the table should not be read as a ranking: those schemes provide conditional privacy, which AUTHBC does not; their group elements are 128 B where our signature is 64 B; and their figures are for a different message. What the comparison does show is that per-message authentication of this kind costs several hundred bytes, against which a telemetry record is small, and that aggregation alone does not change that.

X. LIMITATIONS

Contention is simulated. The hardware experiment has one transmitter. The capacities in Table III come from ns-3 with one collision domain, equal received power and static nodes; capture, hidden terminals, spatial reuse and mobility are absent. Send instants are redrawn within each period. PX4’s own stream, in simulation, holds its period to about ± 0.4 ms,

which is nearer to a strictly periodic source; that source gives the same mean delivery and a spread between runs 6–23 times wider. A contention experiment on two to five radios is the next measurement; the model’s predictions for it are registered in the repository.

Telemetry is synthetic. Record sizes come from a generator. The check against flight logs covers twelve logs of one autopilot at 5 Hz position resolution; at 50 Hz it rests on one simulated flight of the same autopilot.

One PHY, generated loss. Capacity was simulated at 6 Mb/s only. The loss experiment uses independent, Gilbert and length-dependent loss, all generated at the frame level, not a channel trace; under the last, the design is below its target where the baseline’s frame loses 5% (Sec. VIII-B). The stream-signing schemes of Table IX are sized from their definitions and simulated as frames; none is implemented.

Energy and CPU. Only the sender is metered, and only the first format; receiver and radio energy are modelled. The receiver’s CPU share counts verification and hashing, not decoding.

Security scope. Confidentiality, availability, key management and protection against a delaying attacker are outside the design (Sec. III). With one link per frame, continuity inside a frame rests on the frame signature alone.

Post-quantum signatures. An ML-DSA signature is 2420 B [4] and exceeds the 1500 B MTU with any header, so it cannot ride in one frame at any batch size; freshness caps the batch at four, so it could not be amortised either. A post-quantum variant would be a different design, built around fragmentation [4], [38].

Multi-hop. Relayed records and more than one collision domain are not evaluated.

XI. CONCLUSION

For small telemetry records the frame, not the record, is the right unit of authentication, and what a frame must carry decides where it can be sent. A frame that verifies alone needs a header, a chain link, a signature and one self-contained record: eight of the twelve EU863-870 LoRaWAN data rates cannot carry that, at any batch size. Where the frame fits, one signature and one chain link per four records cut the bytes per record by 70.3% and let one 802.11a collision domain carry 124 nodes where per-record signing carries 35. Two constraints shape the design and are easy to get wrong: the batch is capped by freshness, not by the MTU, and a frame that depends on its predecessor does not meet a verifiability target equal to the loss rate. Batching and header design decide what can run; among 64 B signatures the scheme does not. The capacity figures are simulated and await a contention experiment on hardware.

DATA AND CODE AVAILABILITY

Code, result files, simulation scenarios, hardware traces and tests are at <https://github.com/besco1998/fanet-authbc>; every number in this paper is generated from the result files.

ACKNOWLEDGMENT

The author thanks [SUPERVISOR – TO BE COMPLETED] for a review of an earlier version of this paper that changed three of its results.

DECLARATION OF COMPETING INTERESTS

The author declares no competing interests.

USE OF GENERATIVE AI

An AI coding agent assisted with code, analysis and drafting under the author’s direction. The author reviewed the work and takes responsibility for its content.

REFERENCES

- [1] MAVLink Development Team, “MAVLink developer guide: Message signing (authentication),” https://mavlink.io/en/guide/message_signing.html, 2026.
- [2] ETSI, “Intelligent transport systems (ITS); security; security header and certificate formats,” European Telecommunications Standards Institute, Technical Specification TS 103 097, Jun. 2015.
- [3] R. Rotta and P. Mykityn, “Secure multi-hop telemetry broadcasts for UAV swarm communication,” arXiv:2401.11915, 2024.
- [4] G. Twardokus, N. Bindel, H. Rahbari, and S. McCarthy, “When cryptography needs a hand: Practical post-quantum authentication for V2V communications,” in *Network and Distributed System Security Symposium (NDSS)*, 2024.
- [5] R. Gennaro and P. Rohatgi, “How to sign digital streams,” in *Advances in Cryptology — CRYPTO ’97*, ser. LNCS, vol. 1294. Springer, 1997, pp. 180–197.
- [6] C. K. Wong and S. S. Lam, “Digital signatures for flows and multicasts,” *IEEE/ACM Transactions on Networking*, vol. 7, no. 4, pp. 502–513, 1999.
- [7] A. Perrig, R. Canetti, J. D. Tygar, and D. Song, “Efficient authentication and signing of multicast streams over lossy channels,” in *Proc. 2000 IEEE Symposium on Security and Privacy (S&P)*, 2000, pp. 56–73.
- [8] —, “The TESLA broadcast authentication protocol,” *RSA Crypto-Bytes*, vol. 5, no. 2, pp. 2–13, 2002.
- [9] J. Veara, M. Jain, K. Moy, and A. Ranganathan, “TBRD: TESLA authenticated UAS broadcast remote ID,” arXiv:2510.11343, 2025.
- [10] D. Boneh, B. Lynn, and H. Shacham, “Short signatures from the Weil pairing,” in *Advances in Cryptology — ASIACRYPT 2001*, ser. LNCS, vol. 2248. Springer, 2001, pp. 514–532.
- [11] D. Boneh, C. Gentry, B. Lynn, and H. Shacham, “Aggregate and verifiably encrypted signatures from bilinear maps,” in *Advances in Cryptology — EUROCRYPT 2003*, ser. LNCS, vol. 2656. Springer, 2003, pp. 416–432.
- [12] H. Li, C. Shen, H. Huang, and C. Wu, “A certificateless aggregate signature scheme for VANETs with privacy protection properties,” *PLOS ONE*, vol. 20, no. 2, p. e0317047, 2025.
- [13] C. Zhang, R. Lu, X. Lin, P.-H. Ho, and X. Shen, “An efficient identity-based batch verification scheme for vehicular sensor networks,” in *IEEE INFOCOM 2008 — The 27th Conference on Computer Communications*, 2008, pp. 246–250.
- [14] A. S. Rajasekaran, A. Maria, F. Al-Turjman, C. Altrjman, and L. Mostarda, “Anonymous mutual and batch authentication with location privacy of UAV in FANET,” *Drones*, vol. 6, no. 1, p. 14, 2022.
- [15] S. Haber and W. S. Stornetta, “How to time-stamp a digital document,” *Journal of Cryptology*, vol. 3, no. 2, pp. 99–111, 1991.
- [16] P. Mehta, R. Gupta, and S. Tanwar, “Blockchain envisioned UAV networks: Challenges, solutions, and comparisons,” *Computer Communications*, vol. 151, pp. 518–538, 2020.
- [17] C. Gündoğan, C. Amsüss, T. C. Schmidt, and M. Wählisch, “Reliable firmware updates for the information-centric internet of things,” in *Proc. 8th ACM Conf. on Information-Centric Networking (ICN)*. ACM, 2021, pp. 59–70.
- [18] S. Darzi, M. M. Rahman, I. Karim, R. Behnia, A. A. Yavuz, and E. Bertino, “Future-proofing authentication against insecure bootstrapping for 5G networks: Feasibility, resiliency, and accountability,” arXiv:2510.23457, 2025.

- [19] S. Card and R. Moskowitz, "DRIP entity tag (DET) authentication formats and protocols for broadcast remote identification (RID)," IETF, RFC 9575, 2024, A. Wiethuechter, Ed.
- [20] C. Bormann and P. Hoffman, "Concise binary object representation (CBOR)," IETF, RFC 8949, 2020.
- [21] T. Pelkonen, S. Franklin, J. Teller, P. Cavallaro, Q. Huang, J. Meza, and K. Veeraraghavan, "Gorilla: A fast, scalable, in-memory time series database," *Proceedings of the VLDB Endowment*, vol. 8, no. 12, pp. 1816–1827, 2015.
- [22] G. Bianchi, "Performance analysis of the IEEE 802.11 distributed coordination function," *IEEE Journal on Selected Areas in Communications*, vol. 18, no. 3, pp. 535–547, 2000.
- [23] X. Ma and X. Chen, "Saturation performance of IEEE 802.11 broadcast networks," *IEEE Communications Letters*, vol. 11, no. 8, pp. 686–688, Aug. 2007.
- [24] —, "Performance analysis of IEEE 802.11 broadcast scheme in ad hoc wireless LANs," *IEEE Transactions on Vehicular Technology*, vol. 57, no. 6, pp. 3757–3768, Nov. 2008.
- [25] ns-3 Consortium, "ns-3 network simulator, release ns-3.48," <https://www.nsnam.org>, 2026.
- [26] J. Haxhibeqiri, F. Van den Abeele, I. Moerman, and J. Hoebeke, "LoRa scalability: A simulation model based on interference measurements," *Sensors*, vol. 17, no. 6, p. 1193, 2017.
- [27] Q. Zirak, D. Shashev, and S. Shidlovskiy, "Swarm of drones using LoRa flying ad-hoc network," in *2021 International Conference on Information Technology (ICIT)*, 2021, pp. 400–405.
- [28] PX4 Development Team, "PX4-Autopilot: mavlink_main.cpp, configure_streams_to_default," Source code, release v1.17.0, <https://github.com/PX4/PX4-Autopilot>, 2026.
- [29] 3GPP, "Unmanned aerial system (UAS) support in 3GPP," 3rd Generation Partnership Project, Technical Specification TS 22.125, Apr. 2022.
- [30] D. J. Bernstein, N. Duif, T. Lange, P. Schwabe, and B.-Y. Yang, "High-speed high-security signatures," in *Cryptographic Hardware and Embedded Systems — CHES 2011*, ser. LNCS, vol. 6917. Springer, 2011, pp. 124–142.
- [31] National Institute of Standards and Technology, "Digital signature standard (DSS)," NIST, FIPS PUB 186-5, 2023.
- [32] S. Josefsson and I. Liusvaara, "Edwards-curve digital signature algorithm (EdDSA)," IETF, RFC 8032, 2017.
- [33] LoRa Alliance, "RP002-1.0.3 LoRaWAN regional parameters," LoRa Alliance, Technical Specification RP002-1.0.3, 2021.
- [34] Semtech Corporation, "SX1276/77/78/79 137–1020 MHz low power long range transceiver," Semtech Corporation, Datasheet, Mar. 2015.
- [35] D. Boneh, S. Gorbunov, R. S. Wahby, H. Wee, C. A. Wood, and Z. Zhang, "BLS signatures," IRTF Crypto Forum Research Group, Internet-Draft draft-irtf-cfrg-bls-signature-05, Jun. 2022.
- [36] L. Cao, H. Yin, J. Hu, and L. Zhang, "Performance analysis and improvement on DSRC application for V2V communication," arXiv:2102.07023, 2021.
- [37] H. Wang, L. Wang, K. Zhang, J. Li, and Y. Luo, "A conditional privacy-preserving certificateless aggregate signature scheme in the standard model for VANETs," *IEEE Access*, vol. 10, pp. 15 605–15 618, 2022.
- [38] S. Sobati-Moghadam, "Predictive-CSM: Lightweight fragment security for 6LoWPAN IoT networks," arXiv:2506.01767, 2025.